

The background of the image features a dark blue gradient on the left, transitioning into a large, vibrant, abstract shape on the right. This shape is composed of overlapping curved segments in shades of orange, pink, and purple, creating a dynamic, modern aesthetic.

AWS re:Invent

NOV. 27 – DEC. 1, 2023 | LAS VEGAS, NV

B0A315

Dos and don'ts of NoSQL data modeling

Almas Shaikh

Solutions Architect
AWS

Mohammed Fazalullah Qudrath (Faz)

Sr. Developer Advocate
AWS



© 2023, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Any NoSQL DBAs or enthusiasts?



Agenda

SQL vs. NoSQL

Why NoSQL?

NoSQL database types

NoSQL data modeling best practices and anti-patterns

Takeaways

SQL and NoSQL side by side

SQL

NoSQL

Optimized for storage

Optimized for compute

Normalized/relational

Denormalized/hierarchical

Ad hoc queries

Instantiated views

Scale vertically

Scale horizontally

Good for OLAP

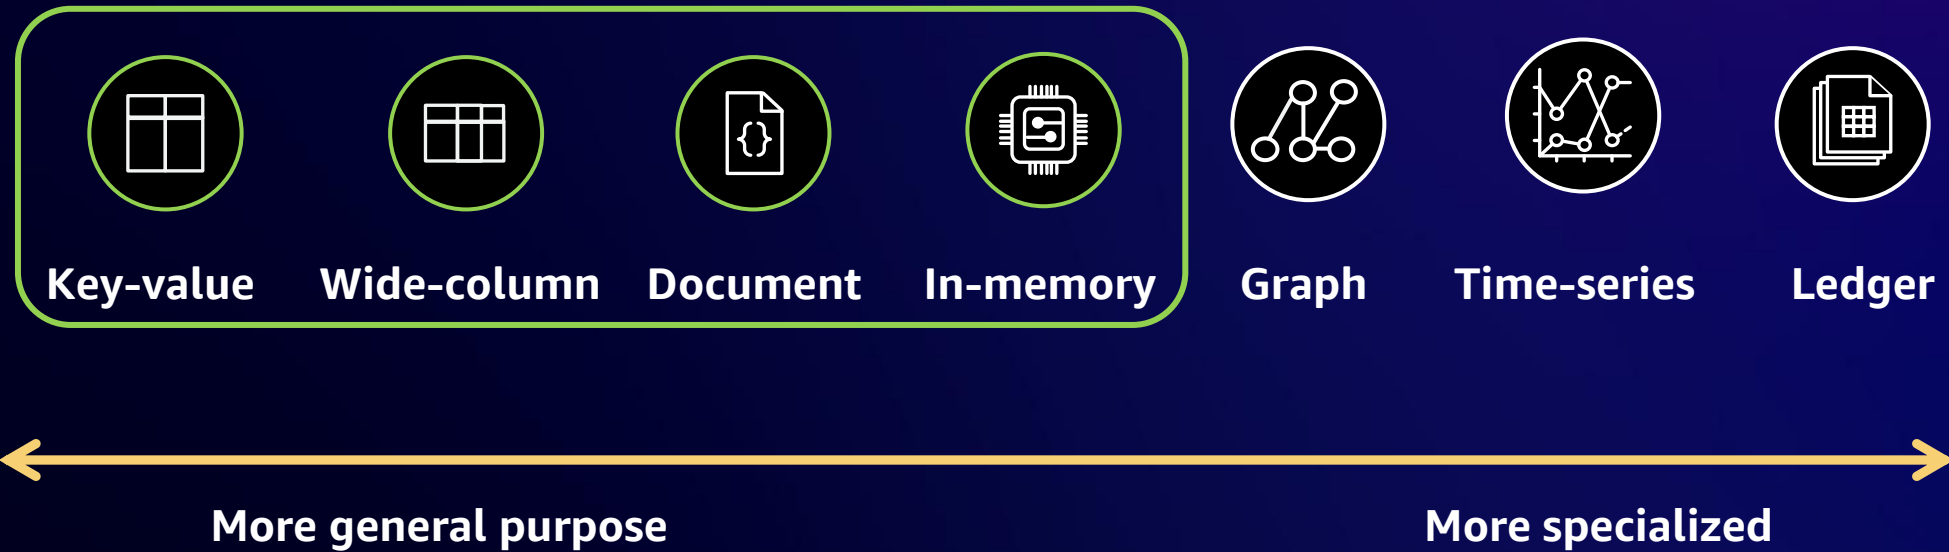
Built for OLTP* at scale

* Most NoSQL databases are built for OLTP workloads

You've moved to
NoSQL because . . .



NoSQL database types



Document databases



Designed for **rich JSON** data models



Supports **complex** access patterns, including indexing and querying nested fields, arrays




Supports **ad-hoc** queries and aggregations

Key-value databases




Designed for **simple**
data models



Key **uniquely**
identifies a record



Joins, GROUP BY and
ORDER BY clauses,
and aggregations
not supported



Provide
extreme scale

Wide-column databases



Designed for **simple**
data models



Key can be
multiple columns

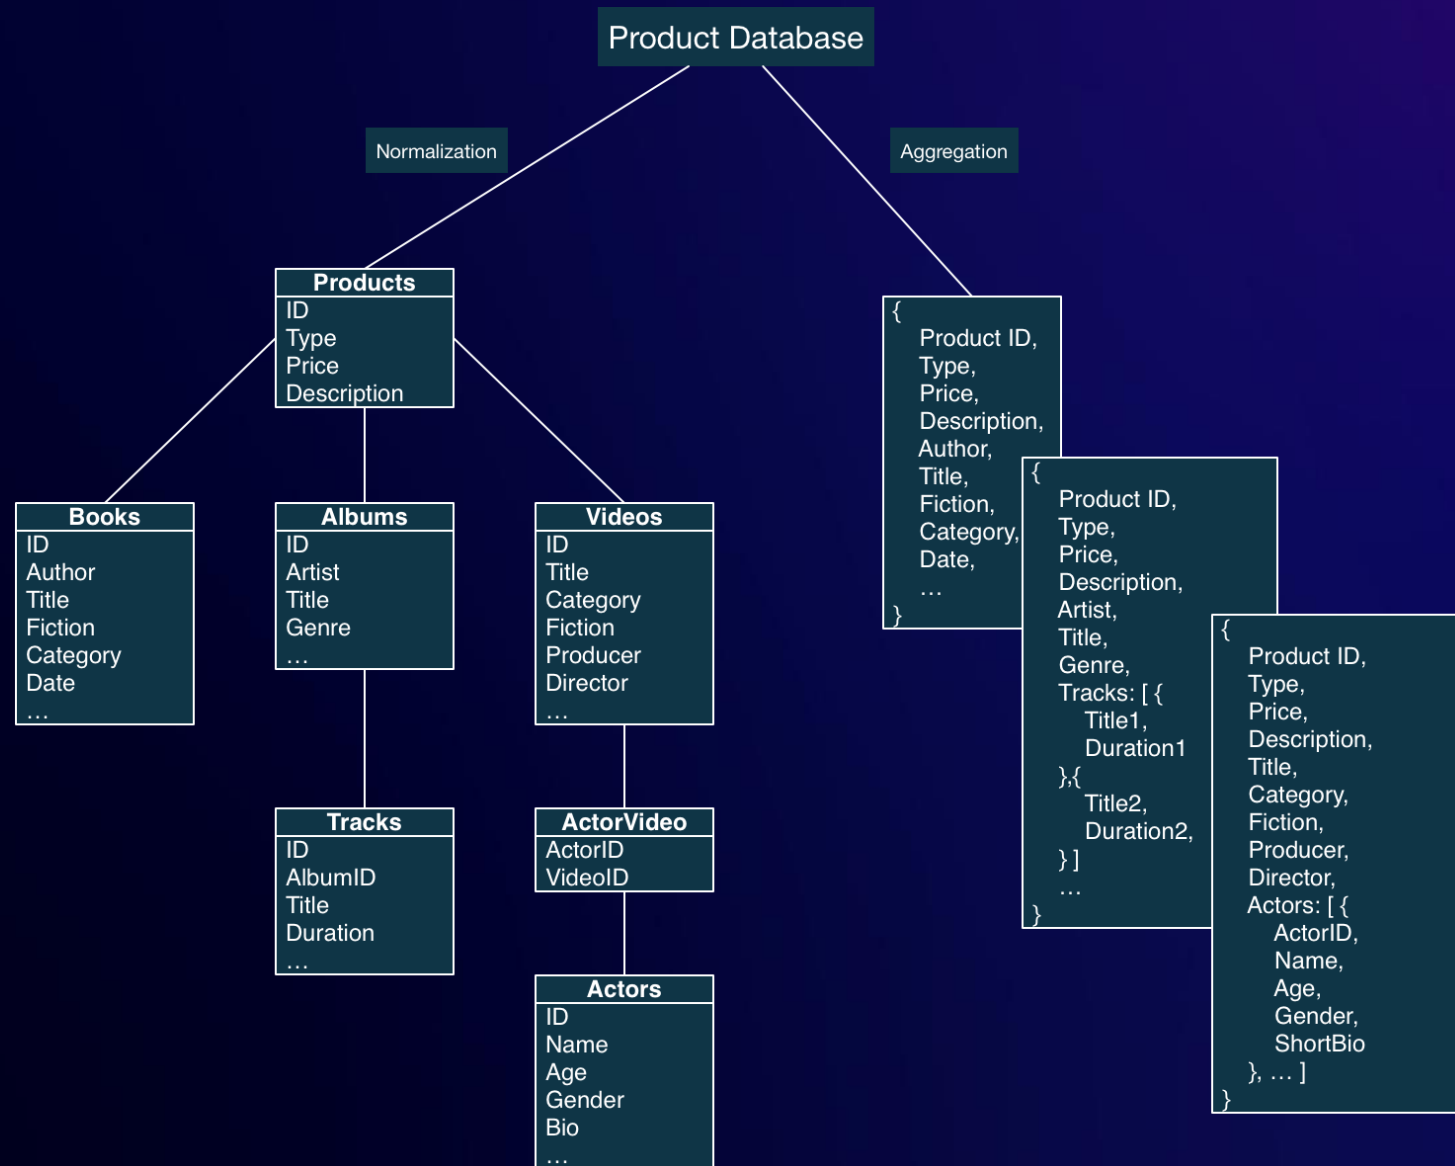


Joins **not supported**
GROUP BY and ORDER
BY and aggregations
limited support



Tables usually have
defined schema

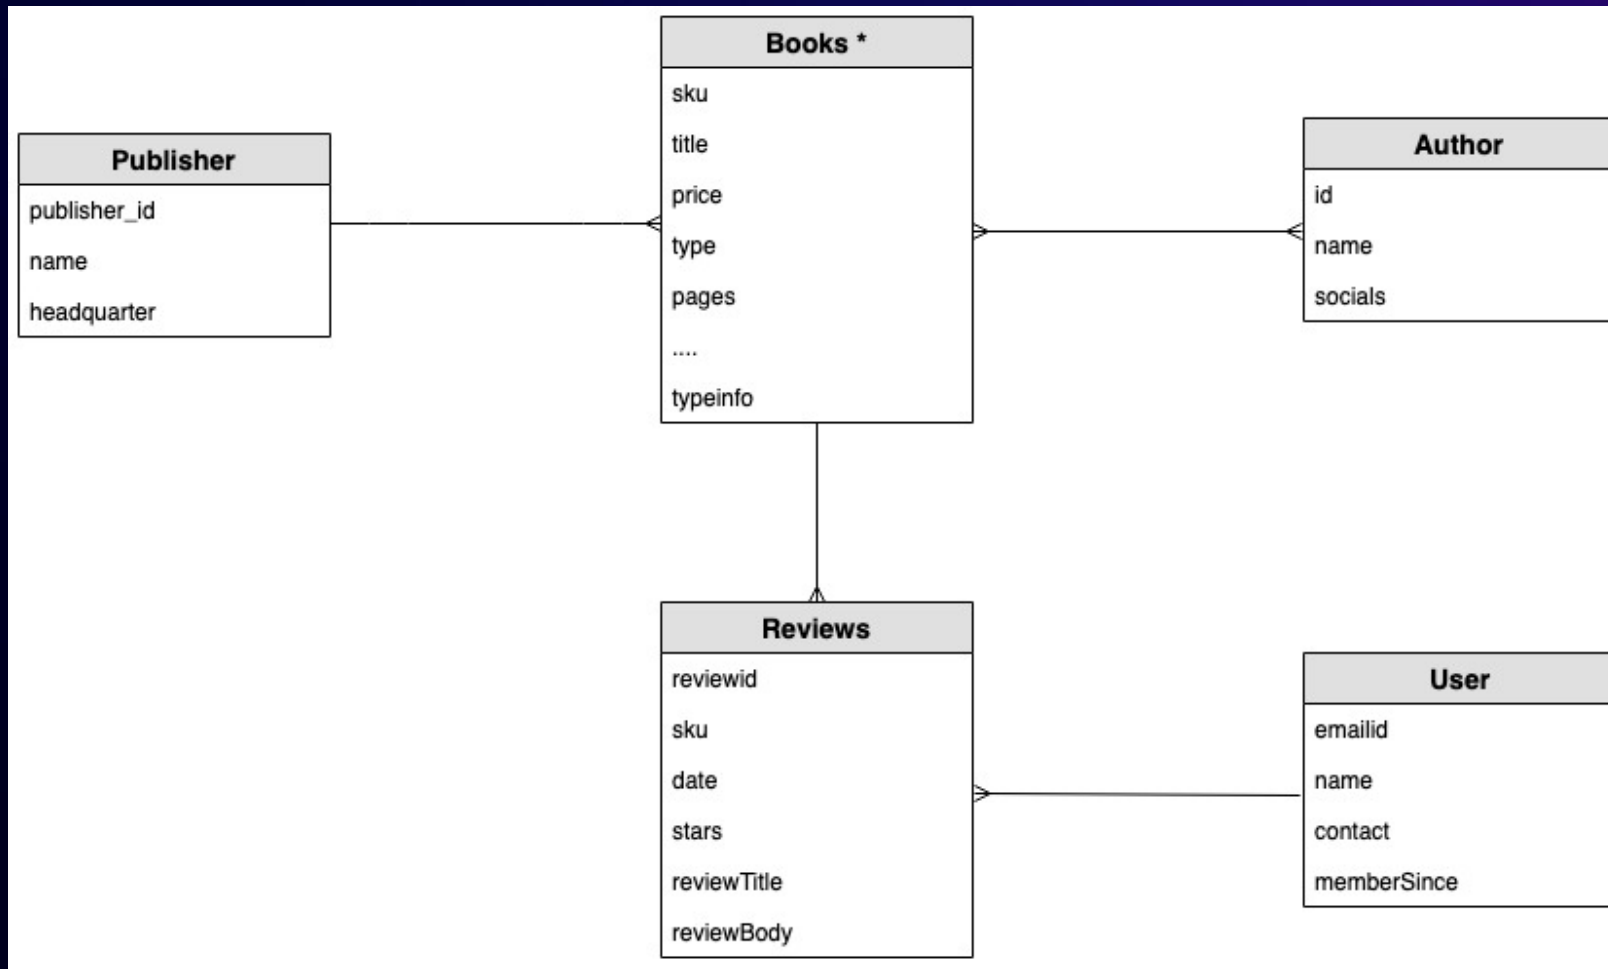
SQL vs. NoSQL design pattern



Data modeling steps

- Define the **use case, entities, and relations**
- Identify the **access** patterns
- Map it to a **model** specific to the database type
- Avoid relational design patterns (**think beyond joins!**)

Book review platform example



Book review platform: Access patterns

- Get **book** info (description and rating)
- Get top **reviews** for a **book**
- Get all **books** by an **author**
- Put/update **book** details
- Add/update **user** details
- Add/update a **review**
- Get **reviews** by **user**

Document database

Constraints:

- BSON/JSON documents
- 16 MB document limit

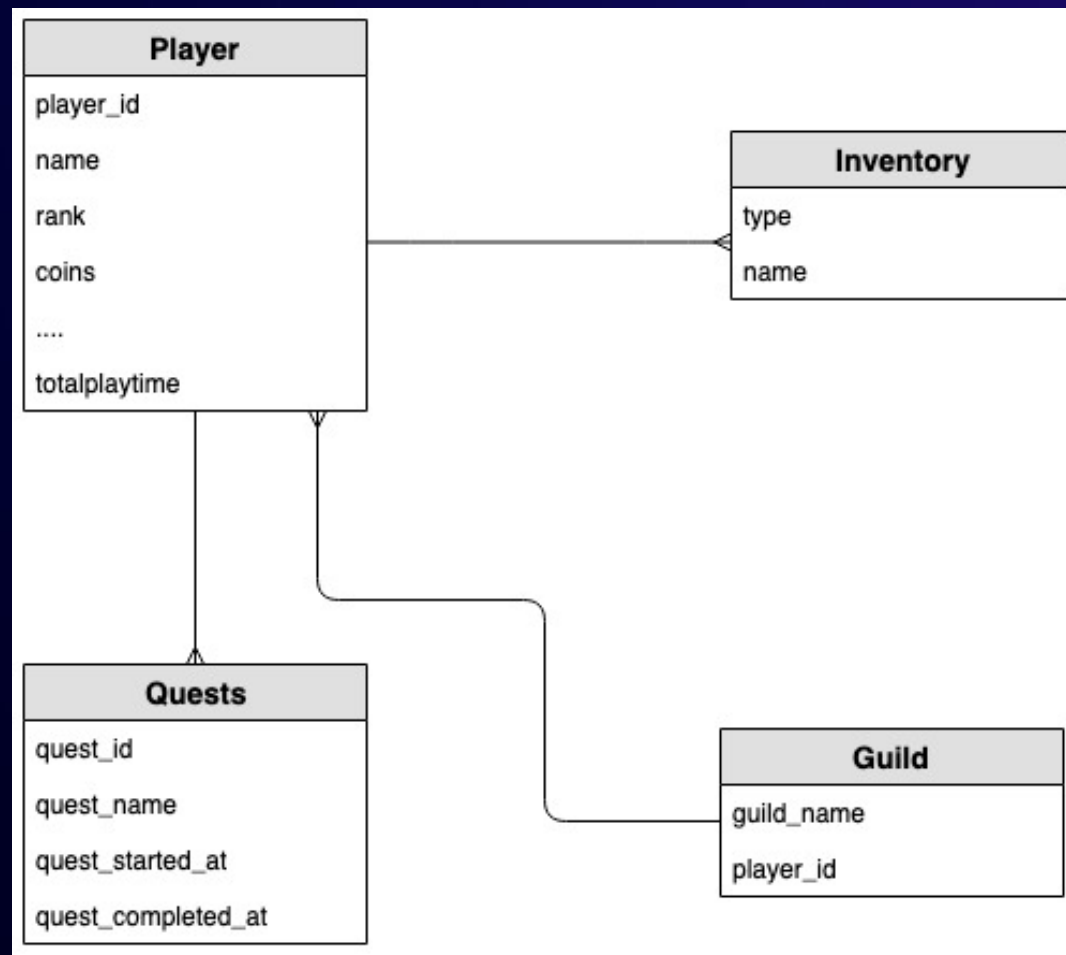
Avoid:

- Bloated documents
- Unbounded arrays
- Diluting the database into too many collections

Learnings:

- Embed unless compelling reason not to
- Polymorphic collections
- Aggregation pipelines – really powerful tool

Gaming example



Gaming: Access patterns

- Create **player**
- Increase coins for **player**
- Add **inventory** to **player**
- Remove **inventory** from **player**
- Add **player** to **guild**
- Get **player**
- Fetch **inventory** for **player**
- Get **players** by **guild**
- Get **quest** details for **player**
- Get **quests** for **player** for last week
- Get completed **quests** for **players**

Key-value database

Constraints:

- 400 KB item limit
- Single attribute partition key
- High cardinality partition key
- LSI limitations

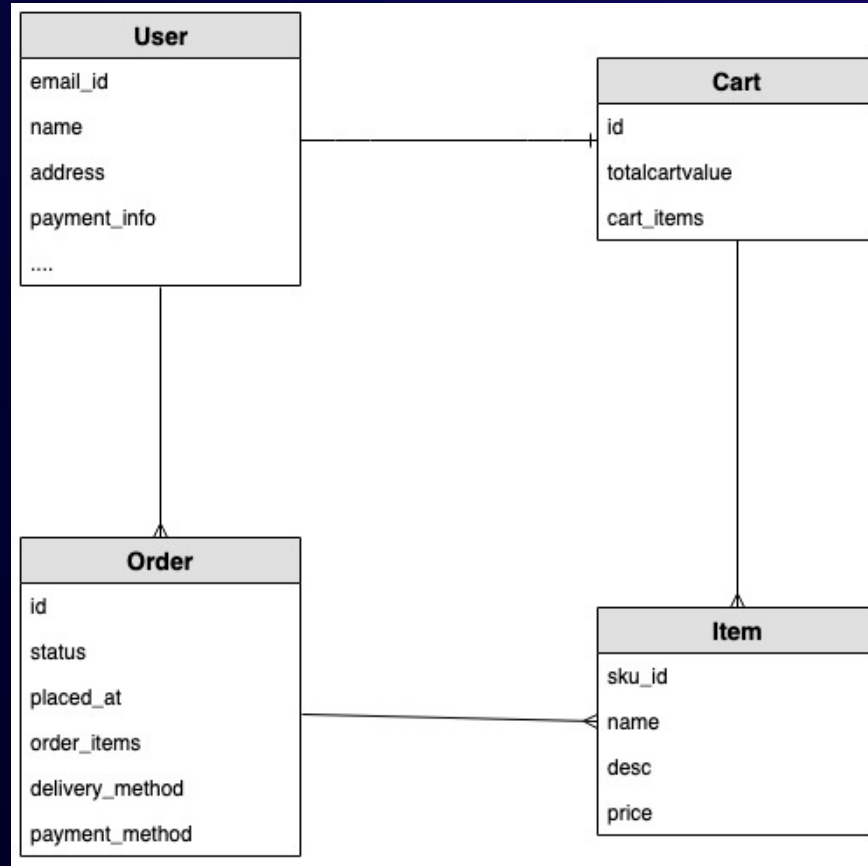
Avoid:

- Large items
- Reads on non-key attributes
- Scans

Learnings:

- Hierarchical items
- Overloading partition key and sort key
- GSIs for additional query capabilities
- Single-table design

Retail cart example



Retail cart: Access patterns

- Get historic details of an **order**
- Create an **order**
- Update **items** in **cart** (add/remove)
- Update status of an **order**
- Get all **orders** placed by **user**
- Display details of an **order**
- Get **orders** placed by a **user** containing an **item**

Wide-column database

Constraints:

- Multi-column partition key
- Multi-column clustering column

Avoid:

- Partitions > 100 MB
- Multi-partition reads/write
- High cardinality indexes

Learnings:

- Access pattern-oriented tables
- Multi-attribute primary key

AWS NoSQL purpose-built databases

CACHING



Amazon
ElastiCache

MEMORY



Amazon
MemoryDB

KEY-VALUE



Amazon
DynamoDB

WIDE-COLUMN



Amazon
Keyspaces

DOCUMENT



Amazon
DocumentDB

TIME-SERIES



Amazon
Timestream

GRAPH



Amazon
Neptune

LEDGER



Amazon
QLDB

From bird's eye view

	Amazon ElastiCache 	Amazon MemoryDB 	Amazon DynamoDB 	Amazon Keyspaces 	Amazon DocumentDB 
Category					
Manageability					
Scale					
API					

Takeaways

- NoSQL data modeling requires a shift in **mindset**
- **De-normalization** is a rule, not an option
- **Work backward** from your access patterns
- Keep re-iterating

Piqued your interest in data modeling?

ATTEND THESE SESSIONS

- **DAT410:** Advanced data modeling with Amazon DynamoDB
- **DAT304:** Data modeling in Amazon DocumentDB (with MongoDB compatibility)

Call to action!

NoSQL Workbench



DocumentDB workshop



DynamoDB workshop



Which NoSQL Database?



Keyspaces workshop



MemoryDB workshop



Thank you!



Please complete the session survey in the mobile app

Almas Shaikh

LinkedIn: [almasshaikh1296](#)

Mohammed Fazalullah Qudrath

LinkedIn: [mohammedfazalullah](#)

